

Data Engineering Technical Assessment: Healthcare Eligibility Pipeline

1. Overview

The goal of this assessment is to build a scalable data pipeline that ingests member eligibility files from multiple healthcare partners, each utilizing different file formats. Your solution should output a unified, standardized dataset ready for downstream consumption.

- **Expected Time Commitment:** 1–2 hours
- **Primary Objective:** Demonstrate the ability to build a **configuration-driven** ingestion process that minimizes code changes when adding new data sources.

2. The Scenario

Our platform receives eligibility data from various partners. Currently, we need to onboard two partners: **Acme Health** and **Better Care**. Their file structures differ significantly. Your pipeline must:

1. Ingest both file formats using a central configuration (avoid hardcoded logic).
2. Transform the raw data into a standardized schema.
3. Output a single, unified dataset.

Sample Data

Partner A: "Acme Health" (`acme.txt`)

- **Format:** Pipe-delimited (|)
- **Data:**

None

```
MBI|FNAME|LNAME|DOB|EMAIL|PHONE  
1234567890A|John|Doe|03/15/1955|JOHN.DOE@EMAIL.COM|5551234567  
9876543210B|Jane|Smith|07/22/1948|jane.smith@email.com|5559876543
```

Partner B: "Better Care" (`bettercare.csv`)

- **Format:** Comma-delimited (,)
- **Data:**

None

```
subscriber_id,first_name,last_name,date_of_birth,email,phone  
BC-001,Alice,Johnson,1965-08-10,alice.j@test.com,555-222-3333
```

BC-002, Charlie, Brown, 1972-03-25, charlie.b@test.com, 5554445555

3. Requirements

A. Configuration-Driven Ingestion

Define partner configurations (e.g., JSON, YAML, or a Dictionary) that specify:

- File delimiters.
- Column mappings (**Partner Column \rightarrow Standard Field**).

Success Metric: Adding a third partner with a different delimiter and column names should only require a configuration update, not a change to the core processing logic.

B. Standardized Output Schema

The final output must adhere to the following transformations:

Field	Transformation
<code>external_id</code>	Map from the partner's unique ID field
<code>first_name</code>	Convert to Title Case
<code>last_name</code>	Convert to Title Case
<code>dob</code>	Format as ISO-8601 (<code>YYYY-MM-DD</code>)
<code>email</code>	Convert to lowercase
<code>phone</code>	Format as <code>XXX-XXX-XXXX</code>

partner_code	Add a hardcoded identifier for the partner
---------------------	--

4. Deliverables

Please provide the following:

1. **Source Code:** A working script/notebook that processes both sample files.
2. **Configuration:** The config files or objects used for both partners.
3. **README:** A brief document explaining:
 - Instructions on how to run the pipeline.
 - A short explanation of how to add a new partner.

Technical Choice

You are free to use the tools you are most comfortable with. We use **Databricks/PySpark** internally, solutions with other tools are perfectly acceptable.

5. Evaluation Rubric

Criteria	What We're Looking For
Functionality	Does the code run and produce the correct standardized output?
Scalability	Is the ingestion truly configuration-driven?
Code Quality	Is the code readable, organized, and maintainable?

Bonus Points (Optional)

- **Validation:** Ensure `external_id` is present; flag or drop rows where it is missing.
- **Error Handling:** Gracefully handle malformed rows or incorrect date formats.